

Symbolic AI

Andre Freitas



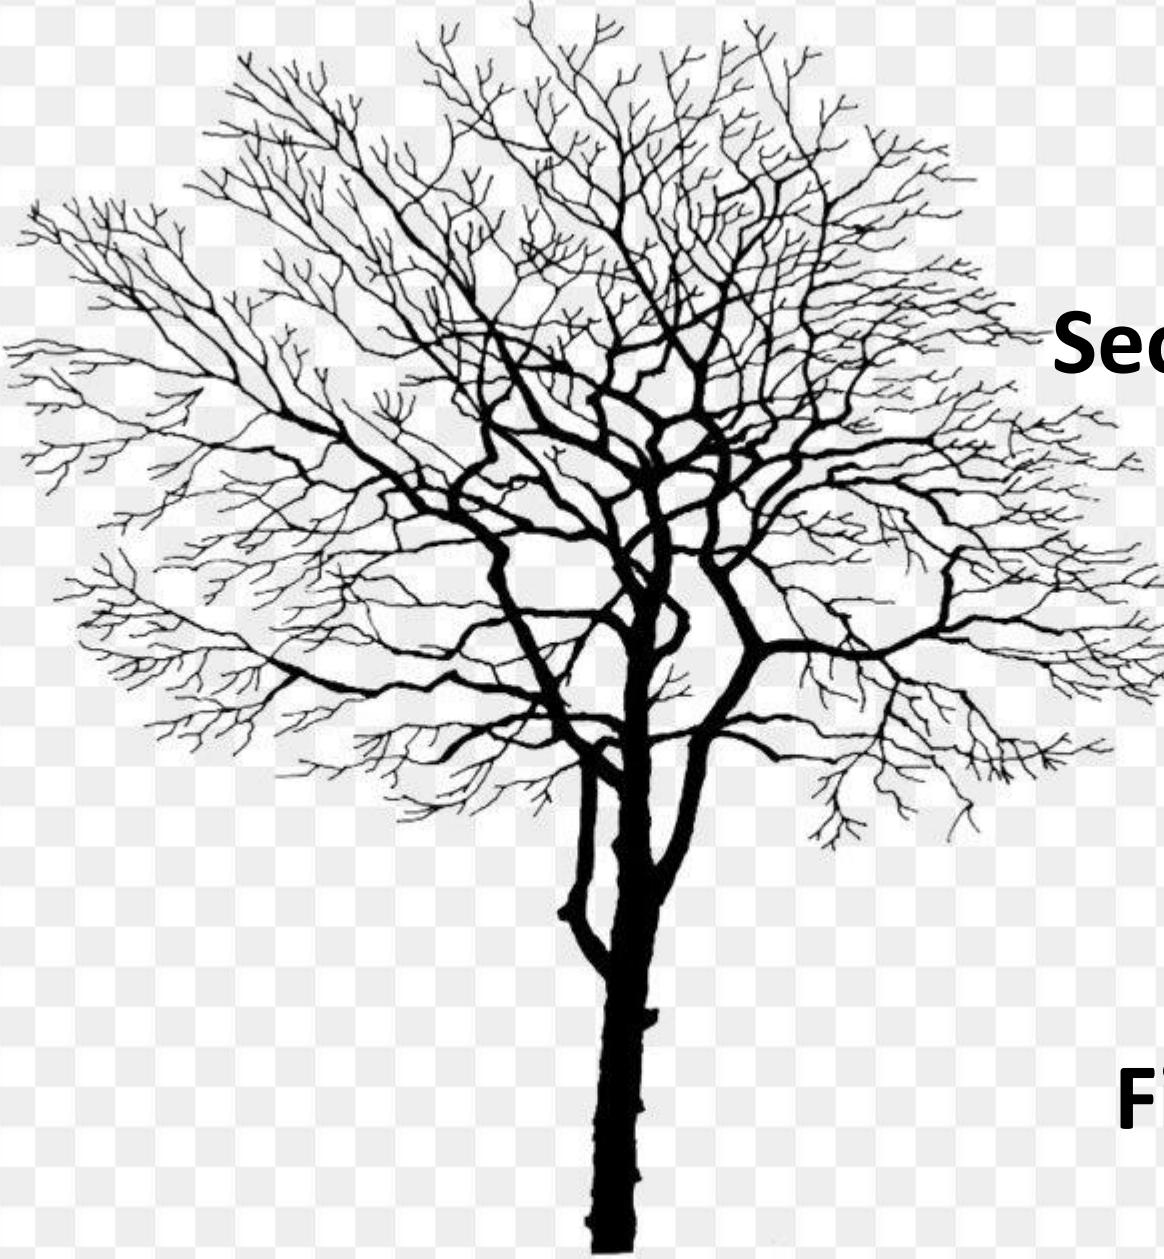
Photo by Vasilyev Alexandr

A close-up photograph of a person's hand gripping a dark-colored, textured gear shift lever in a car. The background shows the interior of the car, including a wooden dashboard and various control buttons.

**Overview of the second
part of the course**

Pedagogical Take

- Giles provided you with the foundations on logical representation, reasoning and programming.
- We will build and expand on it.
- Comparatively, this part of the course will cover more topics (broader strokes).
- Fundamental to provide you with an end-to-end view on Symbolic AI.
- Mastering complexity.



Second Part

First Part

How to Study

- Be in a position to define and explain:
 - The core concepts and algorithms.
 - Why they are relevant?
 - When you should use/not use them?
- Basic application of the core algorithms.

Employability Skills

- This unit is heavily complementary to Machine Learning and Text Mining.
- Contemporary AI is evolving rapidly in the direction of hybrid **neuro-symbolic models**.
- Larger palette to build AI Systems.

[Web](#)[News](#)[Images](#)[Shopping](#)[Videos](#)[More ▾](#)[Search tools](#)

About 186,000,000 results (0.53 seconds)

Michelle Obama (m. 1992)

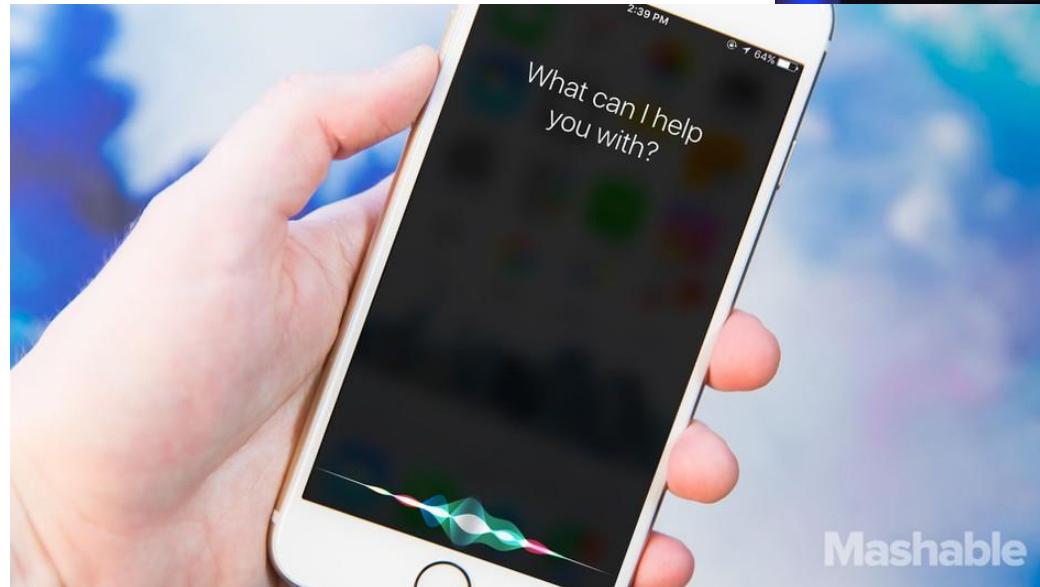
Barack Obama, Spouse



[Michelle Obama - Wikipedia, the free encyclopedia](#)

en.wikipedia.org/wiki/Michelle_Obama ▾ Wikipedia

Michelle LaVaughn Robinson Obama (born January 17, 1964), an American writer, is the wife of the 44th and current President of the United States, Barack Obama. She is also the mother of Malia and Sasha Obama, and the widow of the late Chicago lawyer and U.S. Representative Robert H. "Bobby" Ladd, Jr.



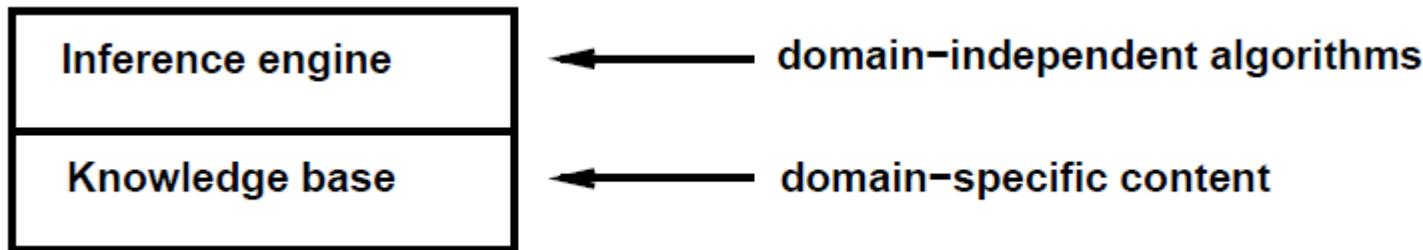
S LARGEST AIRPORT IS NAMED FOR A WORLD WAR II HERO; ITS SECOND LARGEST, FOR A WORLD WAR II BATTLE

What is
Toronto?????

\$36,681

\$

Knowledge Bases



- Knowledge base = set of sentences in a formal language
- Declarative approach to building an agent (or other system):
 - Tell it what it needs to know
- Then it can Ask itself what to do | answers should follow from the KB
- Agents can be viewed at the knowledge level
 - i.e., what they know, regardless of how implemented
- Or at the implementation level
 - i.e., data structures in KB and algorithms that manipulate them

KB Agent

function KB-AGENT(*percept*) **returns** an *action*

static: *KB*, a knowledge base

t, a counter, initially 0, indicating time

 TELL(*KB*, MAKE-PERCEPT-SENTENCE(*percept, t*))

action \leftarrow ASK(*KB*, MAKE-ACTION-QUERY(*t*))

 TELL(*KB*, MAKE-ACTION-SENTENCE(*action, t*))

t \leftarrow *t* + 1

return *action*

A simple knowledge-based agent

- The agent must be able to:
 - Represent states, actions, etc.
 - Incorporate new percepts
 - Update internal representations of the world
 - Deduce hidden properties of the world
 - Deduce appropriate actions
- => sound and complete reasoning with partial information states

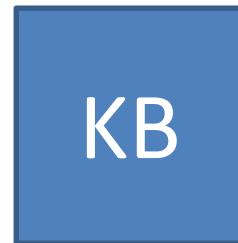
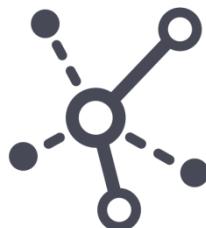
Knowledge in Learning

- Task: to design agents that already know something, and are trying to learn more.
- Agents must have a learning process to gain the background knowledge in the first place
 - Learning taken place afterwards define the agent's incremental/cumulative development
- Agents can start off like normal agents
 - Gain initial knowledge through inductive learning
 - After, uses background knowledge to learn more effectively

Goals of this second part of the course

- Which knowledge representation elements support more expressive, accurate, efficient and general algorithms?
- How to build knowledge bases (under a certain representation)?
- What are the conceptual frameworks and algorithms for knowledge-based learning and inference?

Representation



**Problem
Task**

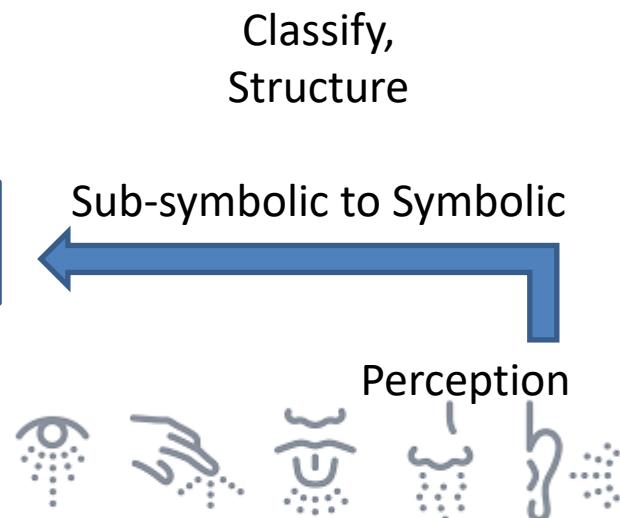


Answer
Explain
Act



Algorithms

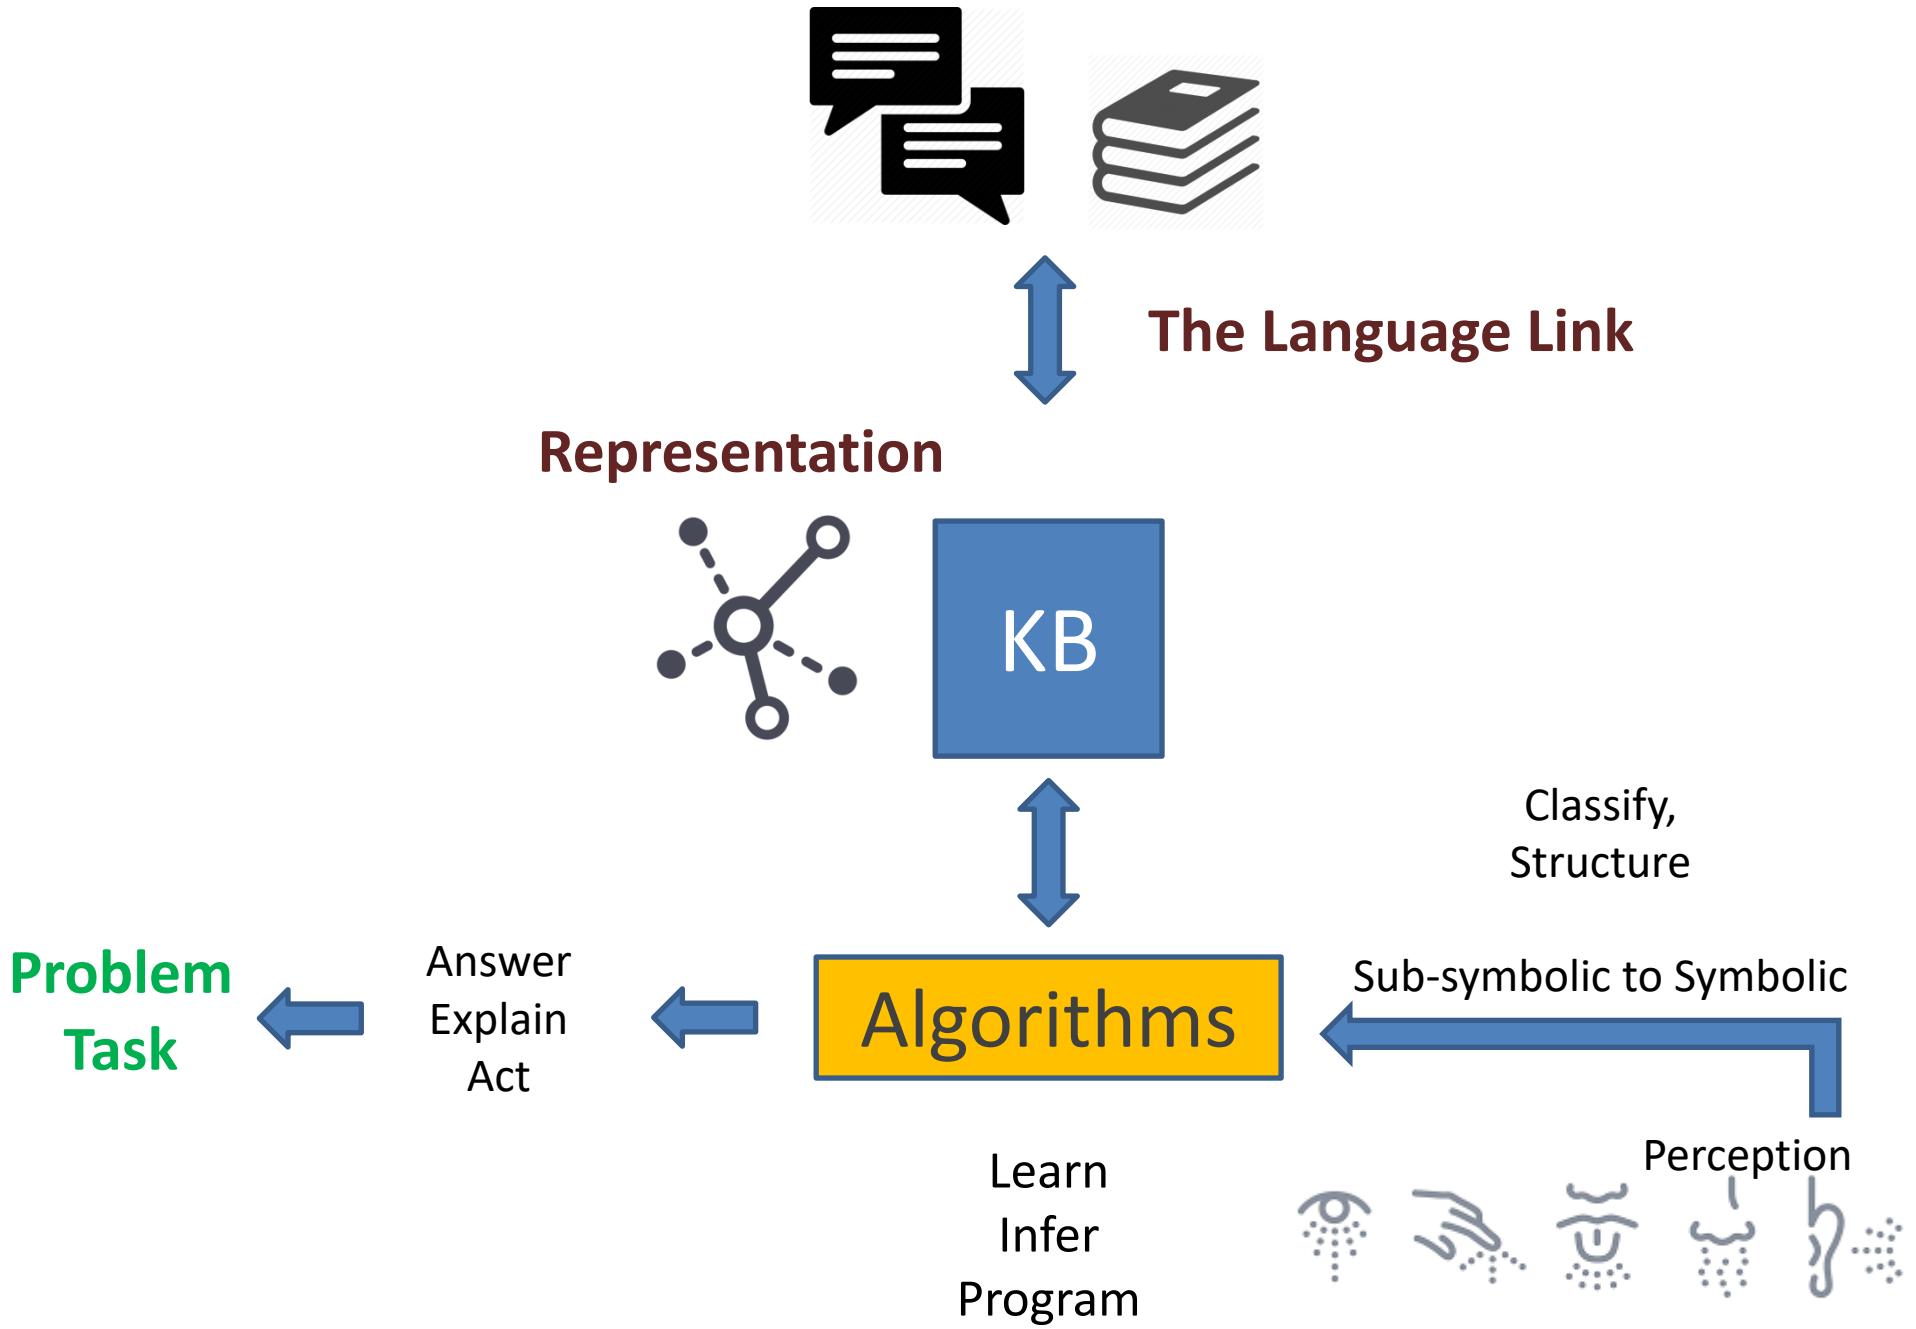
Learn
Infer
Program



Classify,
Structure



Perception

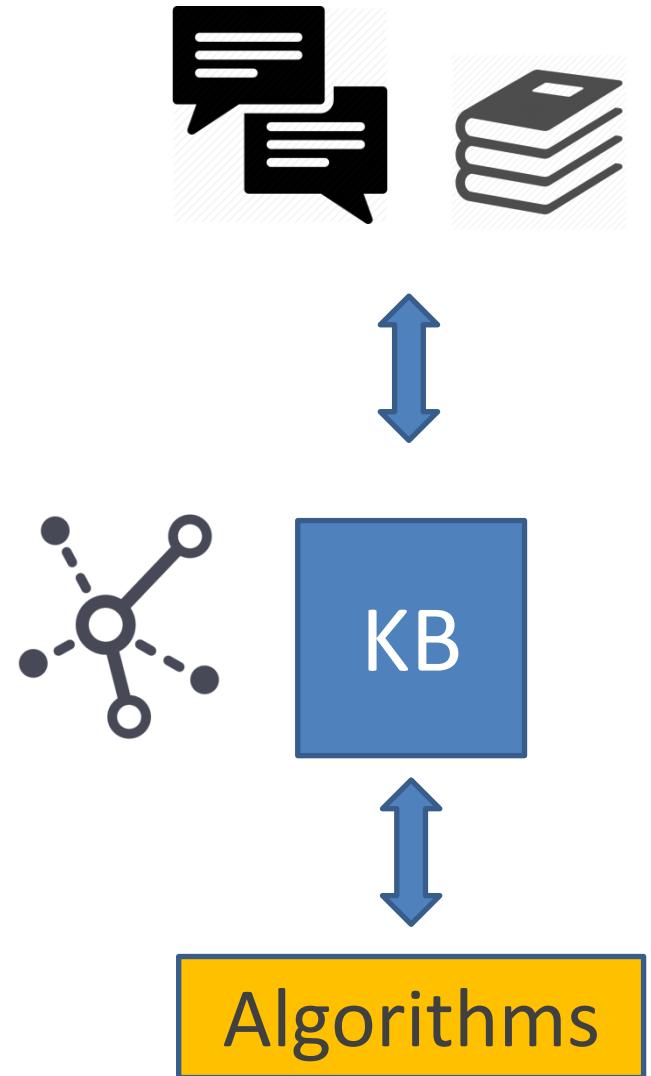


Syllabus

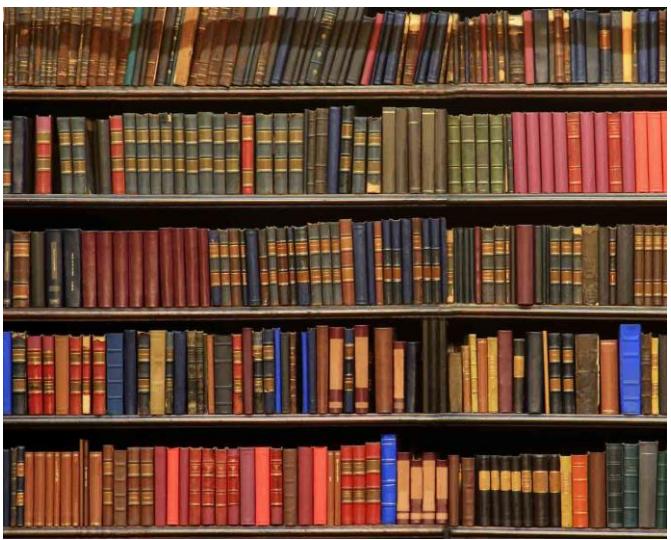
- 2. KB Construction
- 3. Montague Semantics
- 4. Semantic Parsing

- 1. Knowledge Representation

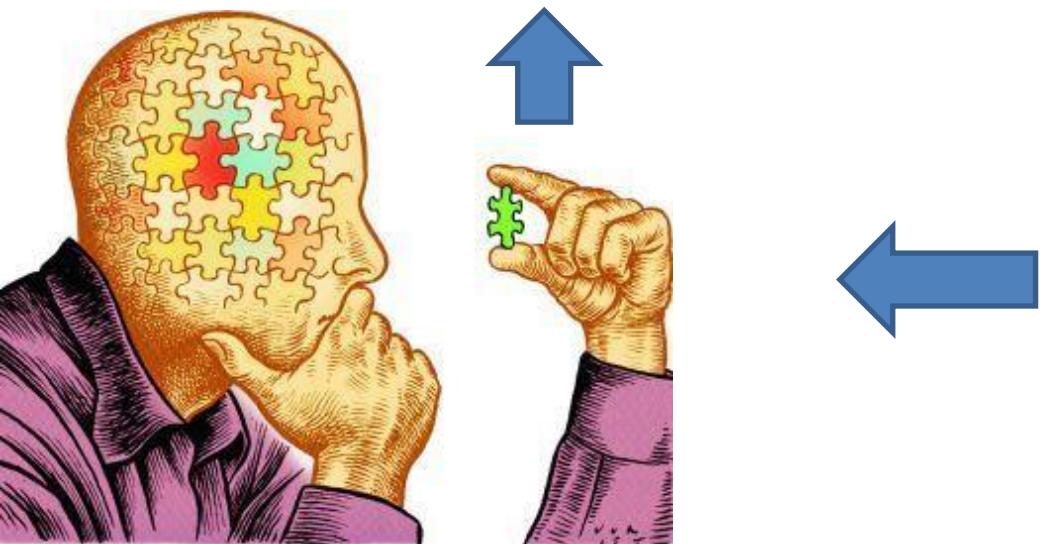
- 5. Explanation-based Learning
- 6. Inductive Logic Programming
- 7. Natural Language Inference
- 8. Neuro-Symbolic Reasoning



Knowledge Representation



Communication of the Representation



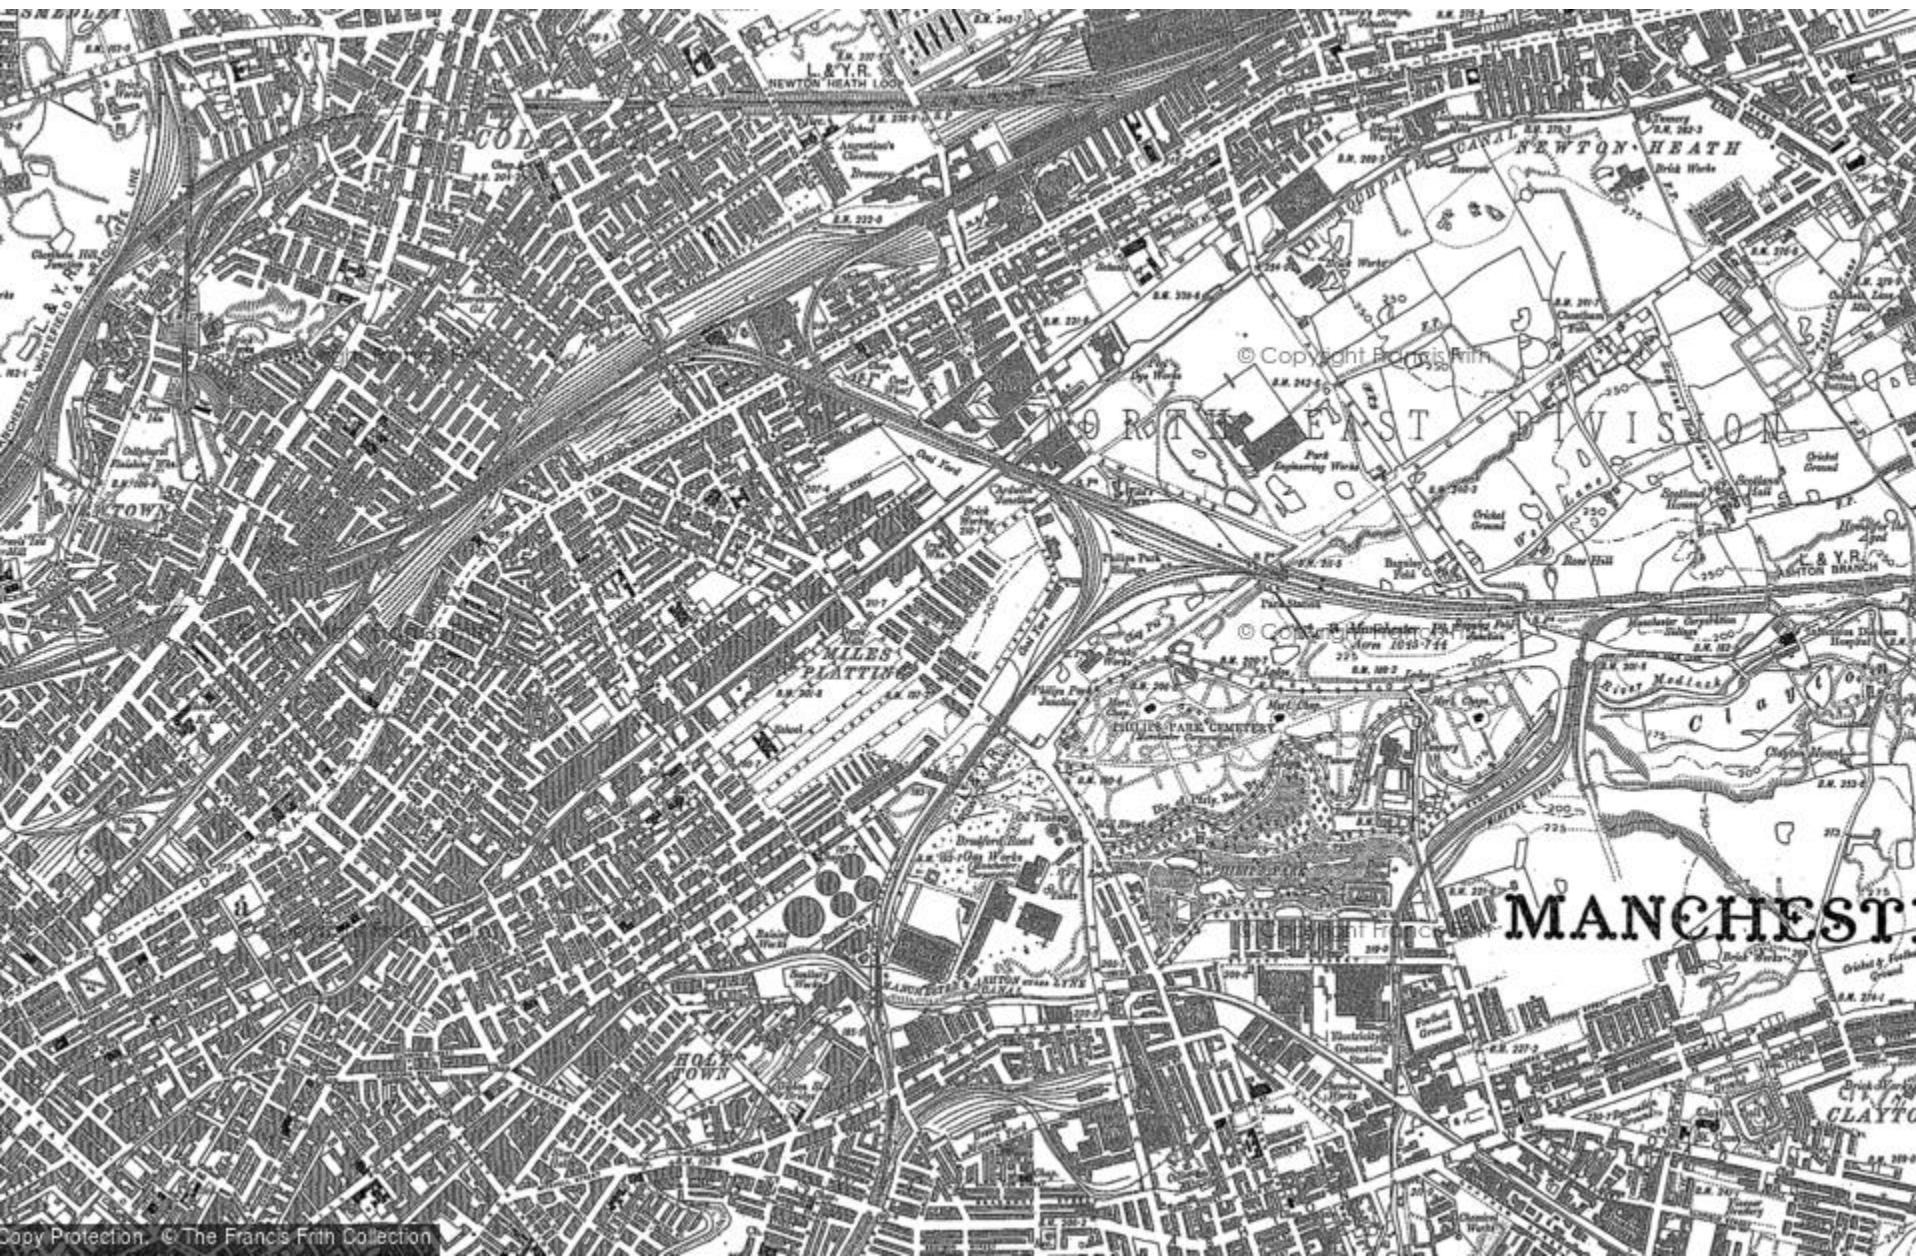
Representation
of the Reality

Structure of the
Reality





Contains Ordnance Survey data © Crown copyright and database right 2011





Services

- 1 Altrincham – Bury
- 2 Altrincham – Piccadilly
- 3 Ashton-under-Lyne – Eccles
- 4 Bury – Piccadilly
- 5 East Didsbury – Rochdale Town Centre
- 6 Manchester Airport – Victoria*
- 7 MediaCityUK – Etihad Campus

Manchester Airport – Victoria*

MediaCityUK – Etihad Campus

* Early services operate on a 20 minute frequency between Manchester Airport and Didsbury. Please check journey planning posters or [tfgm.com](#) before travelling.

Key

- Metrolink stop
- Bus interchange
- Rail interchange
- Change with other services
- Park + Ride with number of spaces
- Car park fewer than 50 spaces
- Cycle Hub membership required



Ceci n'est pas une pipe.

“Human knowledge is a process of approximation. In the focus of experience, there is comparative clarity. But the discrimination of this clarity leads into the penumbral background. There are always questions left over. The problem is to discriminate exactly what we know vaguely.”

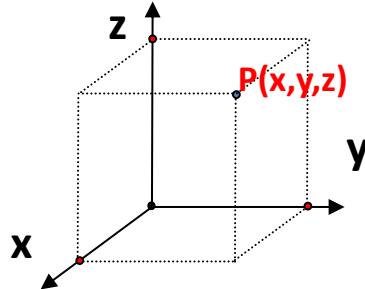
Alfred North Whitehead

KR: Five Roles

- 1. Surrogate
 - That is, a representation
- 2. Expression of ontological commitment
 - of the world
- 3. Theory of intelligent reasoning
 - and our knowledge of it
- 4. Medium of efficient computation
 - that is accessible to programs
- 5. Medium of human expression
 - and usable

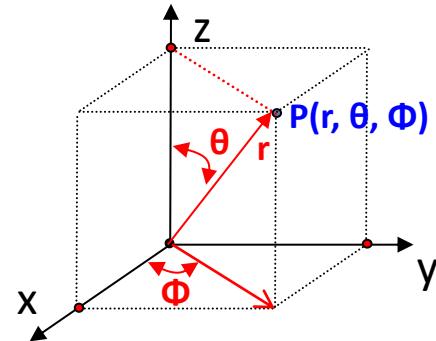
Cartesian Coordinates

$P(x, y, z)$



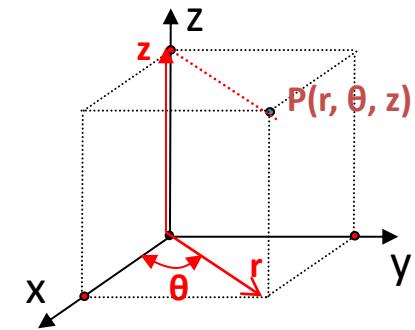
Spherical Coordinates

$P(r, \theta, \Phi)$



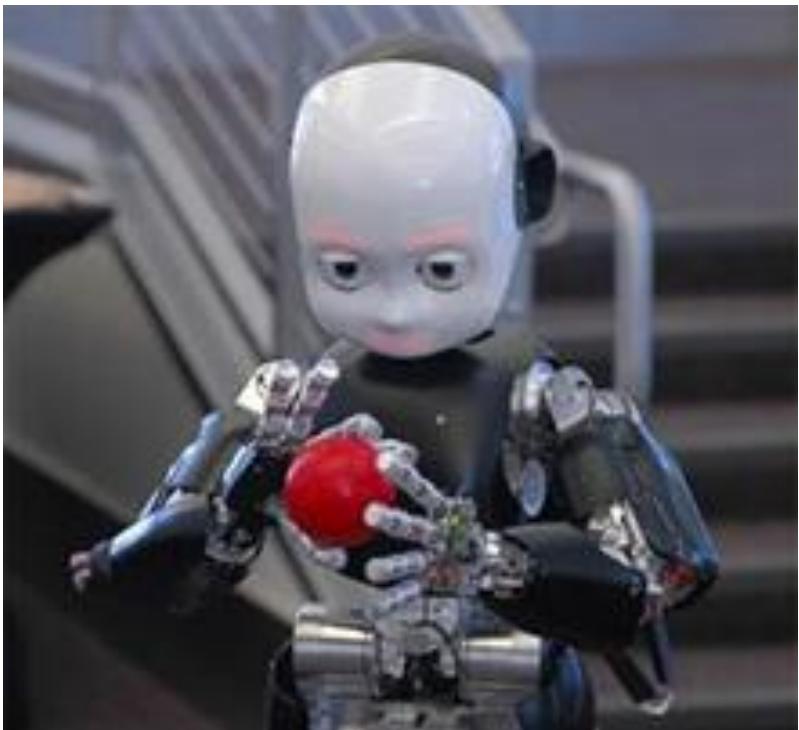
Cylindrical Coordinates

$P(r, \theta, z)$



$$ds^2 = \left(1 - \frac{2m}{r}\right) dt^2 - \frac{1}{\left(1 - \frac{2m}{r}\right)} dr^2 - (r)^2(d\theta^2 + \sin^2(\theta)d\phi^2)$$

Representations deeply impact on learning and inference



Embodied representations

An **apple** is a sweet, edible [fruit](#) produced by an **apple tree** (*Malus pumila*). Apple [trees](#) are [cultivated](#) worldwide and are the most widely grown species in the [genus Malus](#). The tree originated in [Central Asia](#), where its wild ancestor, *Malus sieversii*, is still found today. Apples have been grown for thousands of years in [Asia](#) and [Europe](#) and were brought to North America by [European colonists](#). Apples have religious and [mythological](#) significance in many cultures, including [Norse](#), [Greek](#) and [European Christian traditions](#).

Symbolic representations

“The distinctive feature of brains such as the one we own is their uncanny ability to create maps...

But when brains make maps, they are also creating images, the main currency of our minds. Ultimately consciousness allows us to experience maps as images, to manipulate those images, and to apply reasoning to them.”

Antonio Damasio (2010)

Semantics

=

Formal meaning representation
model (lots of data)

+

inference model

This behaves a lot
like intelligence!

> 2000 years of tradition!

Semantics

=

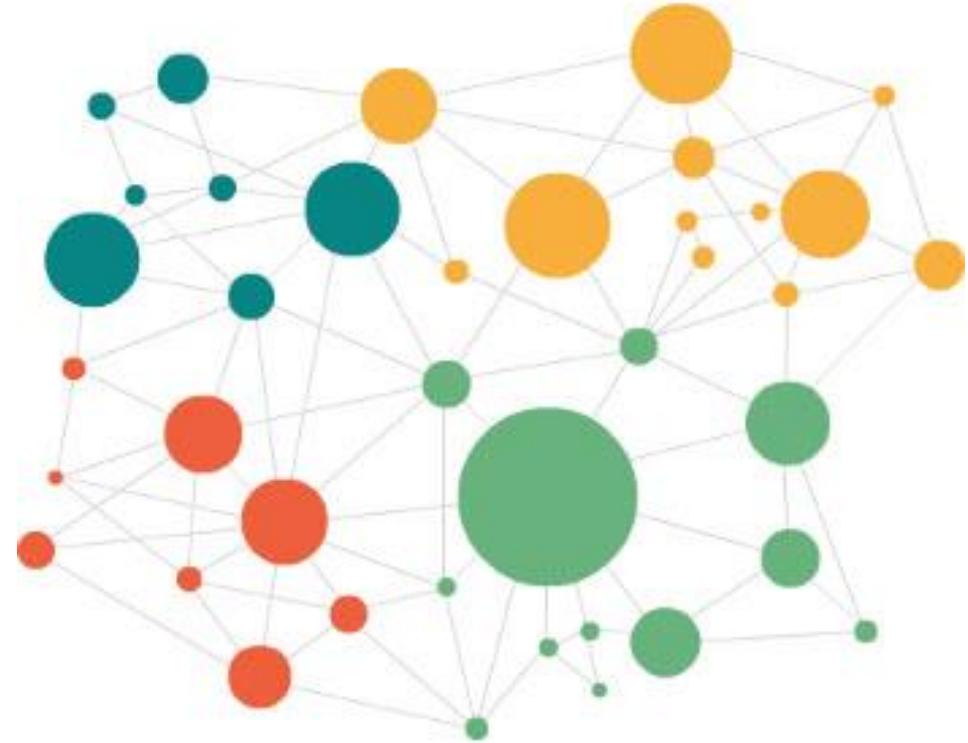
Logics, linguistics, philosophy, cognitive sciences, computer science

Formal meaning representation model (lots of data)

+

inference model

This behaves a lot like intelligence!



Building Knowledge Bases

Data

Intelligence

Structure/Semantics



Unstructured
Data

KB Construction

Structured
Data

Easy to generate

Easy to analyze
(computationally)

Consistent
Comparable
Processable

From Text to Structure

...

Barack Obama went with his daughter Malia to the baseball game.

...

...

Today, during an official visit, Natasha called to her father, the president of the United States.

...

From Text to Structure

...
Barack Obama went with
his daughter **Malia** to the
baseball game.

...
Today, during an official
visit, **Natasha** called to
her father, the president
of the United States.

From Text to Structure

...
Barack Obama went with
his daughter **Malia** to the
baseball game.
...

...
Today, during an official
visit, **Natasha** called to
her father, the **president
of the United States**.
...

Co-reference resolution

From Text to Structure

...
Barack Obama went with
his daughter **Malia** to the
baseball game.
...

...
Today, during an official
visit, **Natasha** called to
her father, the **president**
of the United States.
...

Regularities in Natural Language

Barack Obama went with his daughter Malia to the baseball game.

Barack/NNP

Obama/NNP

went/VBD

with/IN

his/PRP\$

daughter/NN

Malia/NN

to/TO

the/DT

baseball/NN

game/NN

./.

Regularities in Natural Language

Barack Obama went with his daughter Malia to the baseball game.

Barack/**NNP**



Proper noun

Obama/**NNP**

went/**VBD**

with/**IN**

his/**PRP\$**

daughter/**NN**

Malia/**NN**

to/**TO**

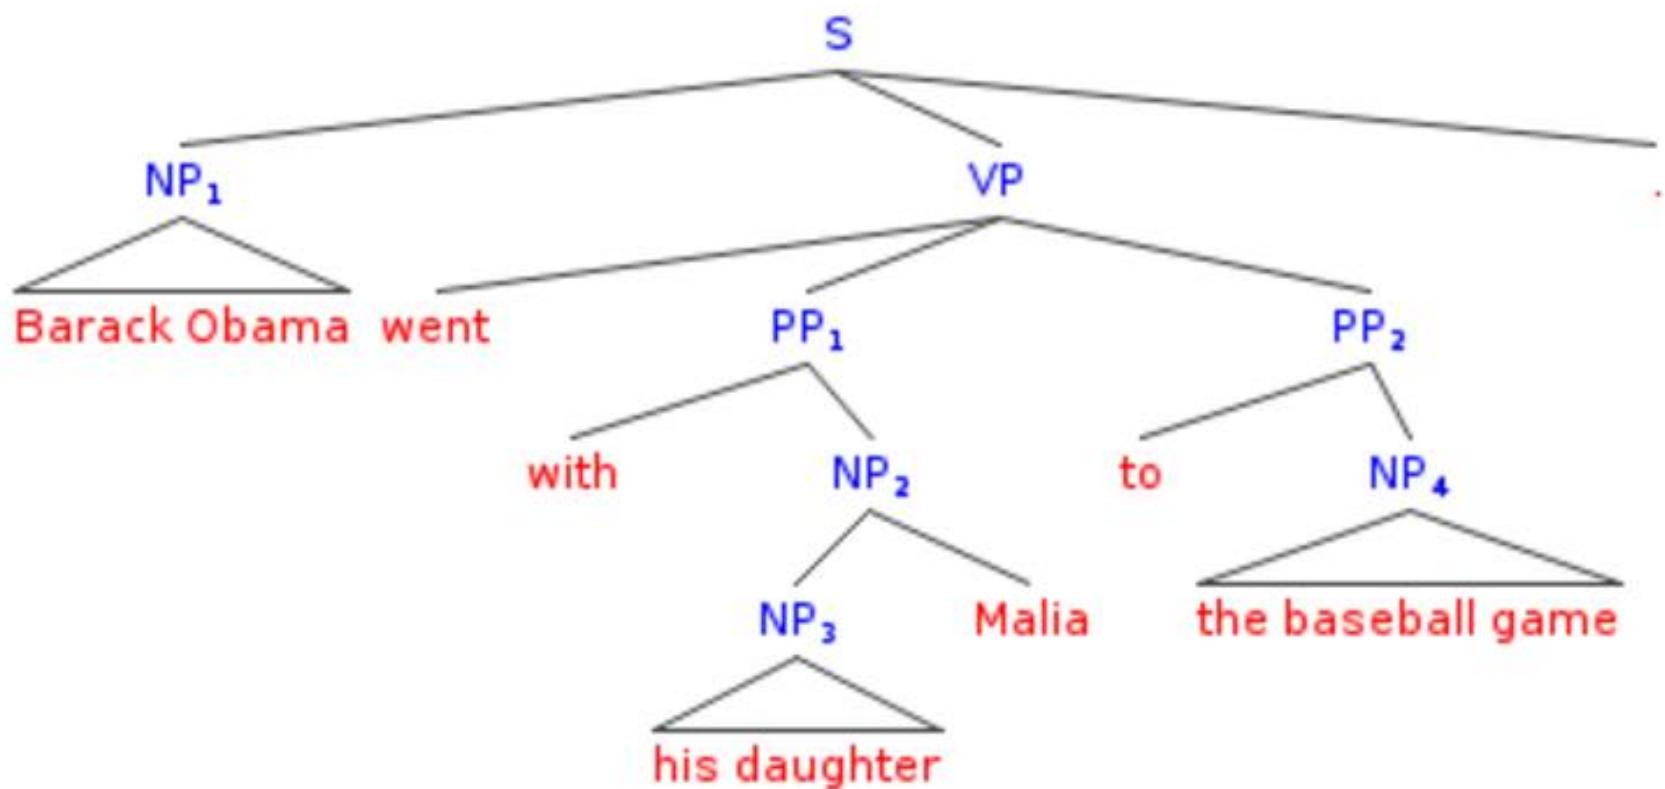
the/**DT**

baseball/**NN**

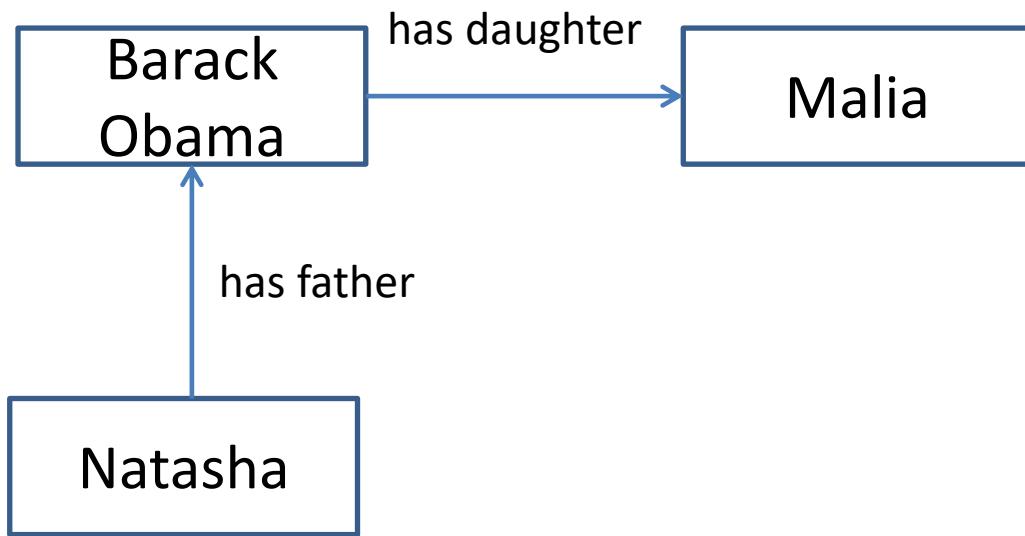
game/**NN**

./.
./.

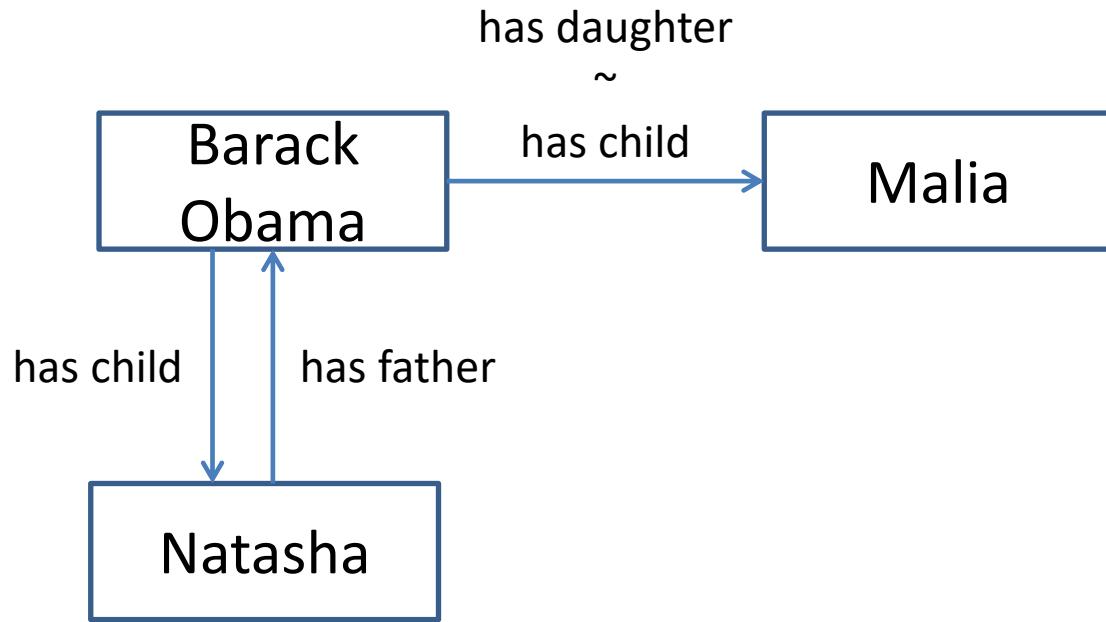
Regularities in Natural Language



Structural/Logical Form



Structural/Logical Form



Applying some logical or corpus-based inference we get

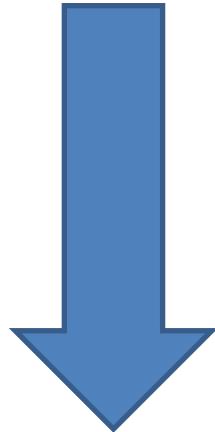
Rephrasing it

- has_child(Barack_Obama, Malia)
- has_child(Barack_Obama, Natasha)
- ...

Now we can answer this query

- *How many children does Barack Obama have?*

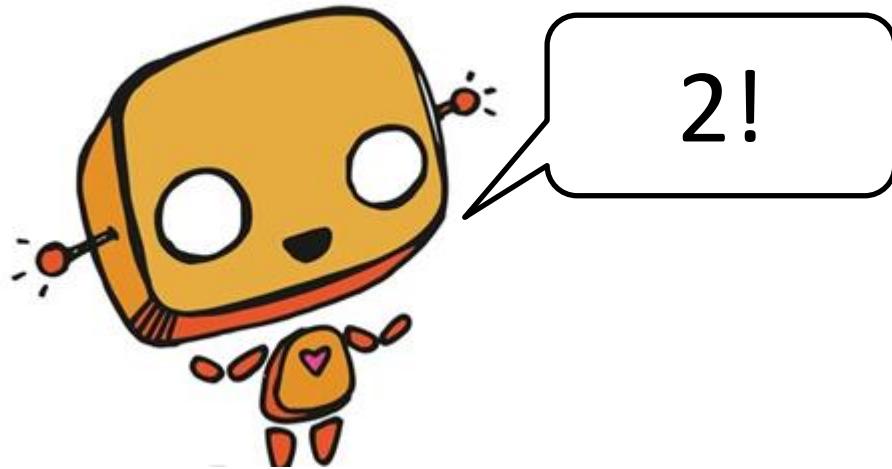
some magic called **semantic parsing** goes on ...



- `count(has_child(Barack_Obama, ?x))`

It computes!

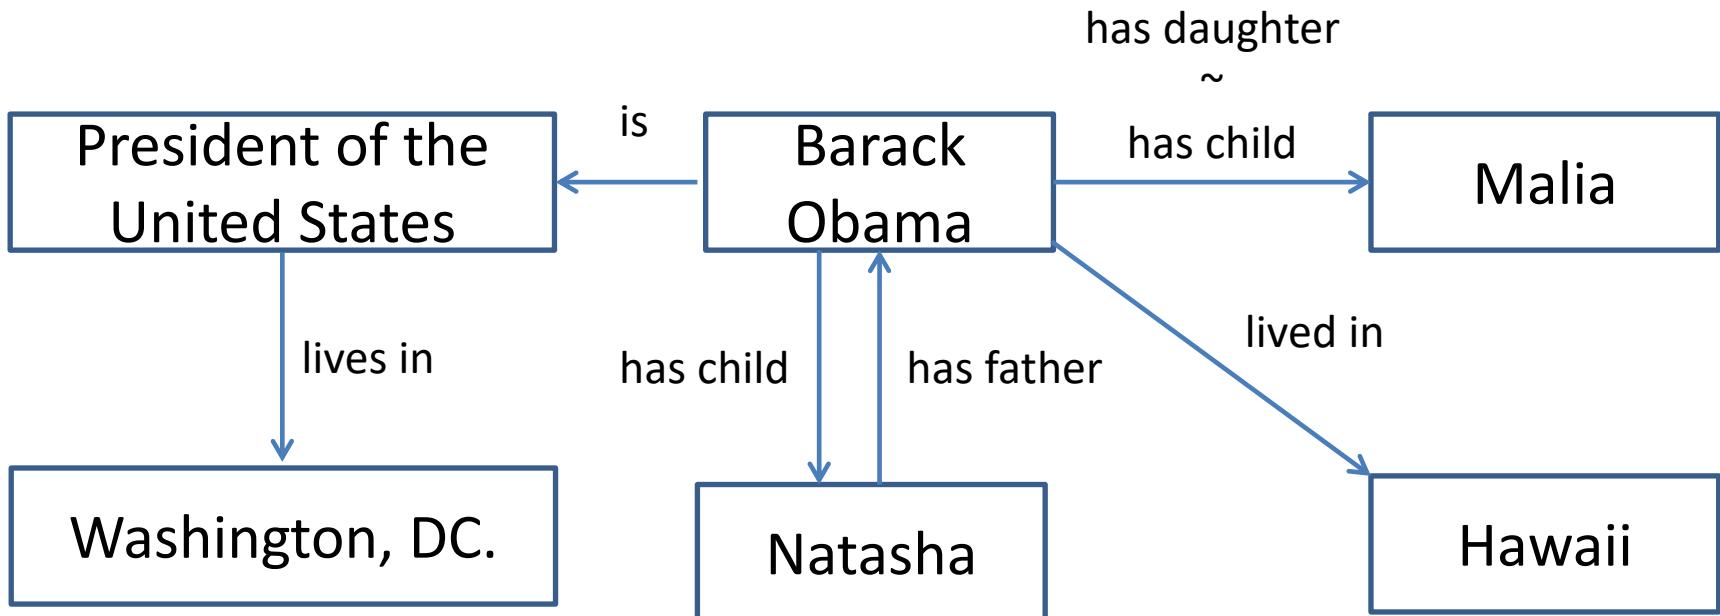
Query: count(has_child(Barack_Obama, ?x))



KB:

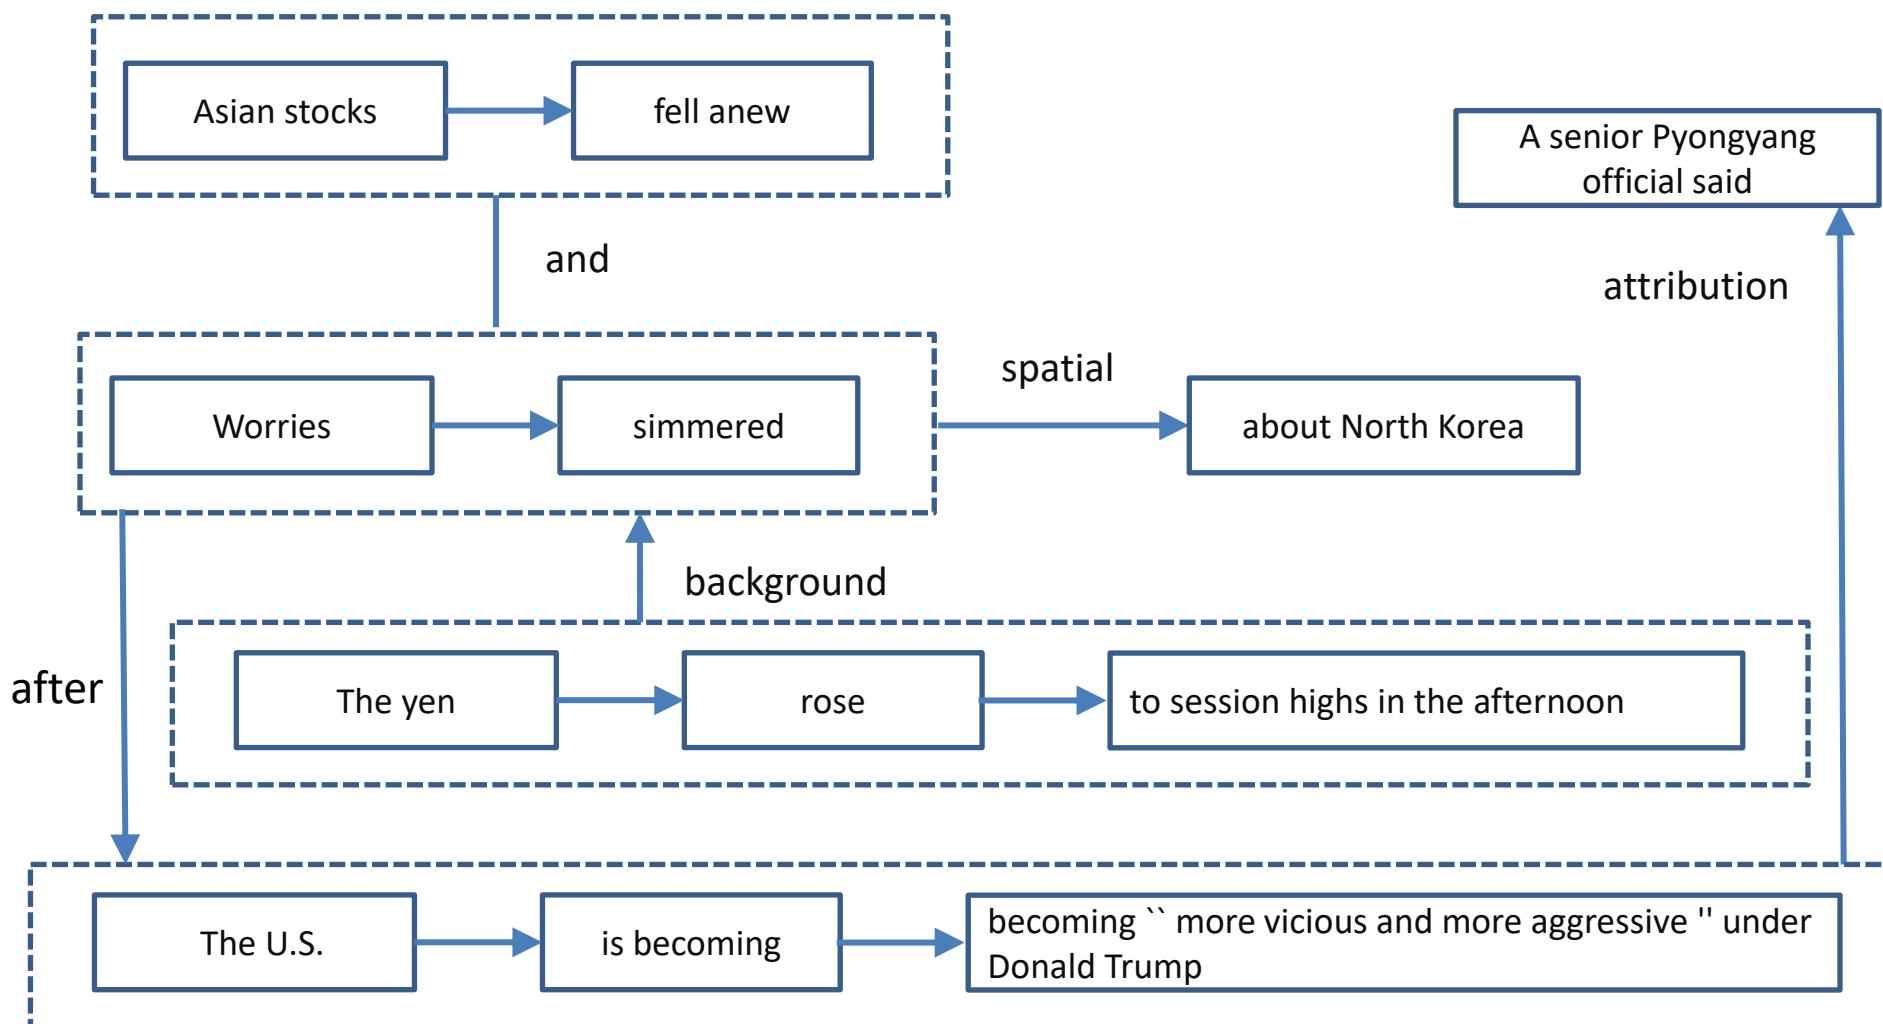
- has_child(Barack_Obama, Malia)
- has_child(Barack_Obama, Natasha)
- ...

Extrapolating



Semantic representation

Asian stocks fell anew and the yen rose to session highs in the afternoon as worries about North Korea simmered, after a senior Pyongyang official said the U.S. is becoming ``more vicious and more aggressive'' under President Donald Trump .



Semantic Parsing using CCGs

$$\frac{\text{show} \quad \text{me} \quad \frac{\text{flights}}{\lambda x. flight(x)} \quad \frac{\text{to}}{\lambda y. \lambda x. to(x, y)} \quad \frac{\text{Boston}}{BOSTON}}{\lambda x. to(x, BOSTON)} \rightarrow$$
$$\frac{\lambda f. \lambda x. f(x) \wedge to(x, BOSTON)}{\lambda x. flight(x) \wedge to(x, BOSTON)} \leftarrow$$
$$\frac{S}{\lambda x. flight(x) \wedge to(x, BOSTON)} \rightarrow$$

Symbol

Tree

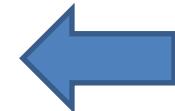


Any cognitive representation for long, vertical, usually green with a wood basis

Tree is the name of a set

Tree is a noun

Some Goal



Similarity, discrimination



Representation

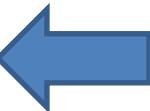
Structure of the Reality

Symbol



Any cognitive representation
for long,
vertical,
usually green
with a wood basis

**Some
Goal**



Similarity,
discrimination

Extension
of the set



Representation

Structure of the Reality

Symbol

Operating on the Representation

Tree



Any cognitive representation for long, vertical, usually green with a wood basis

The tallest tree

Some Goal



Similarity, discrimination



Representation

Structure of the Reality

Symbol

Tree



Any cognitive representation for long, vertical, usually green with a wood basis

Operating on the Representation

Definite article: “get me one”

The tallest tree

Some Goal
Similarity, discrimination



Representation

Structure of the Reality

Symbol

Tree



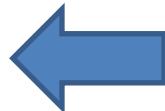
Any cognitive representation for long, vertical, usually green with a wood basis

Operating on the Representation

Superlative adjective : "top most"

The tallest tree

Some Goal



Similarity, discrimination



Structure of the Reality

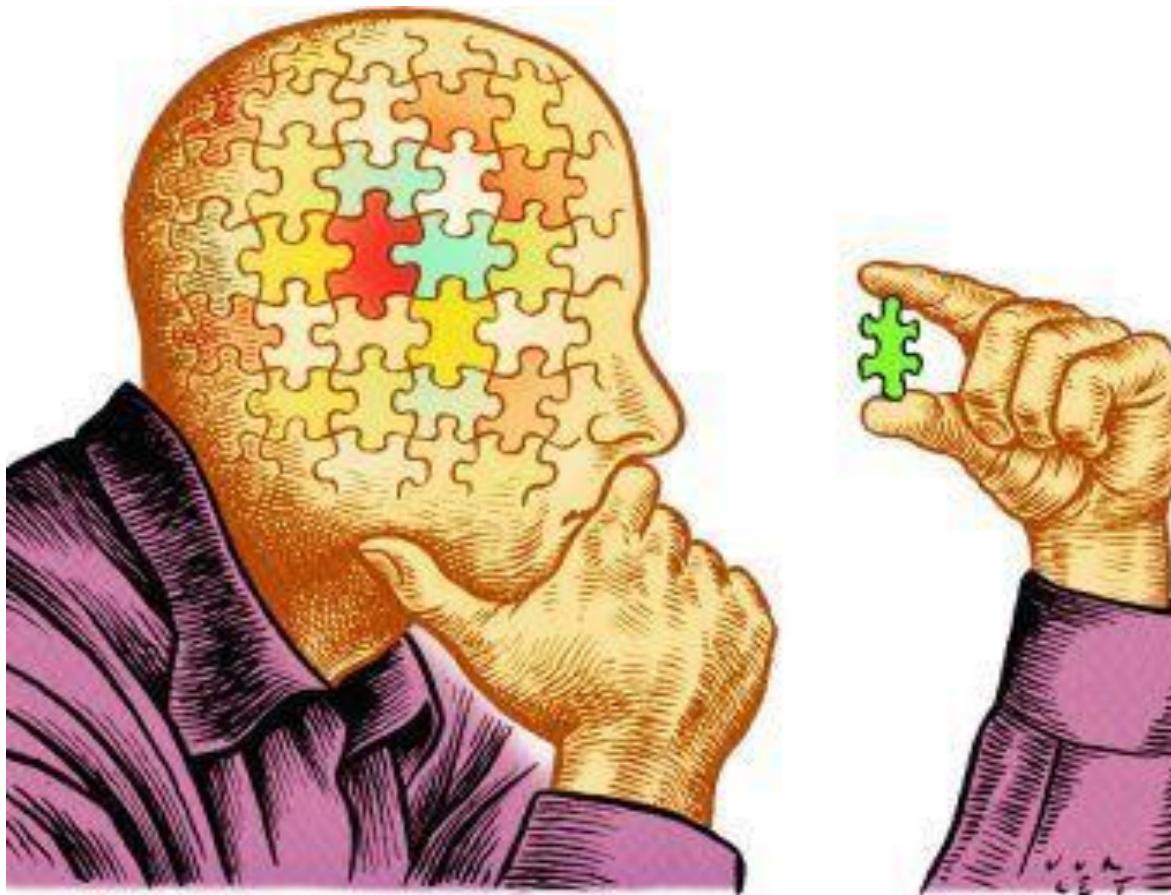
Natural Language Representation

- There is a mapping between natural language and knowledge representation.
- Looking at natural language is looking at the representation (constrained by the communication medium).

Language as a ‘Geological outcrop’ of our cognitive representation.



https://www.geocaching.com/geocache/GC3YMK9_one-day-geology-of-oman-1?guid=c3516272-9eca-4c10-ae14-8dabd9346b98



Knowledge in Learning and Inference

Inductive Logic Programming

- ILP algorithms are constructive induction algorithms
 - Able to create new predicates to facilitate the expression of explanatory hypotheses

- Express Grandparent
 - Empty background
 - Hypotheses are long and complicated

$\text{Grandparent}(x, y) \Leftrightarrow$

$$\begin{aligned} & [\exists z \text{ Mother}(x, z) \wedge \text{Mother}(z, y)] \\ \vee & [\exists z \text{ Mother}(x, z) \wedge \text{Father}(z, y)] \\ \vee & [\exists z \text{ Father}(x, z) \wedge \text{Mother}(z, y)] \\ \vee & [\exists z \text{ Father}(x, z) \wedge \text{Father}(z, y)] \end{aligned}$$

Chu

Inductive Logic Programming

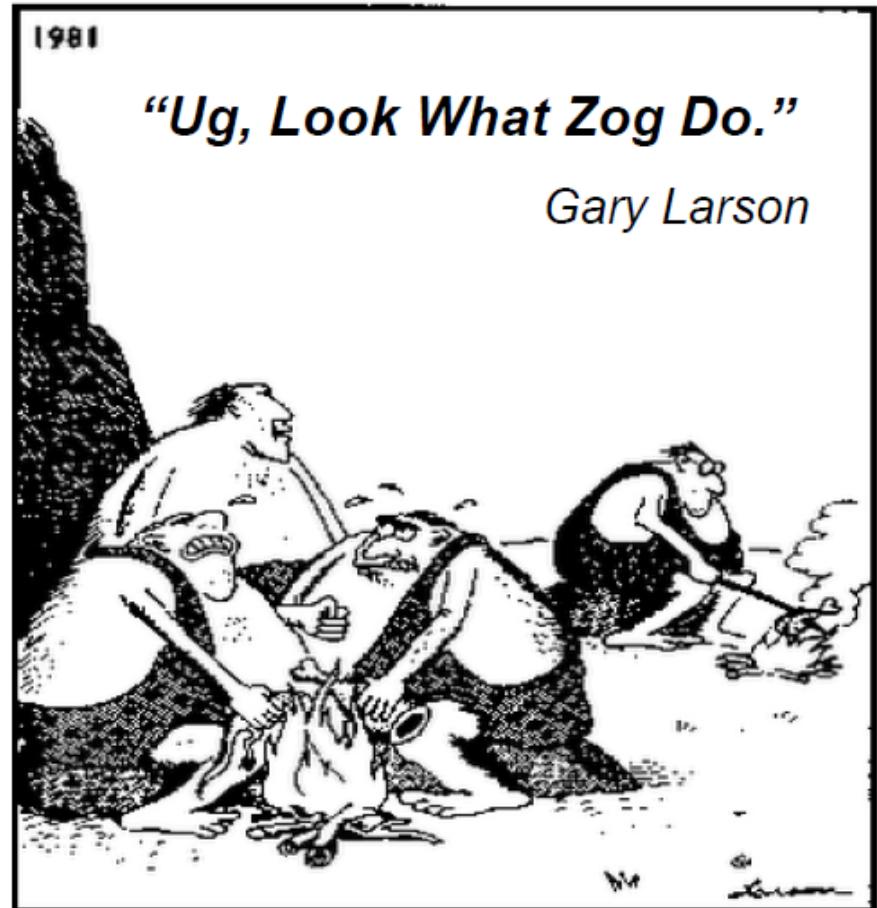
- By creating a new predicate, the definition of Grandparent can be reduced

$$\text{Parent}(x, y) \Leftrightarrow [\text{Mother}(x, y) \vee \text{Father}(x, y)]$$
$$\text{Grandparent}(x, y) \Leftrightarrow [\exists z \text{ Parent}(x, z) \wedge \text{Parent}(z, y)]$$

- Background knowledge can reduce the size of hypotheses required to explain the observations

Explanation-based Learning

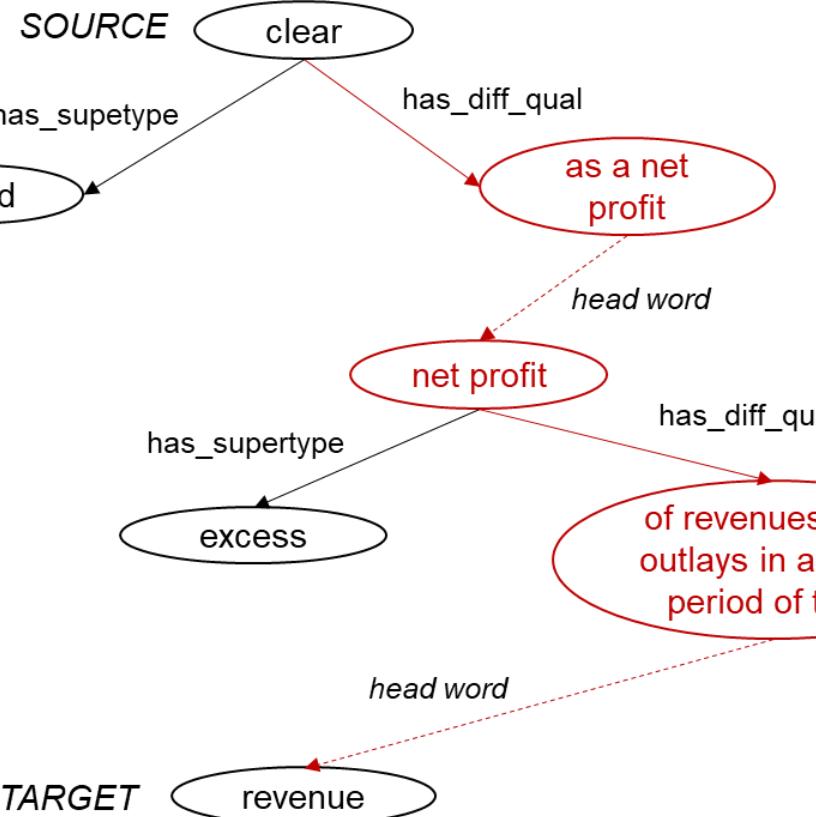
- **Explanation-based Learning (EBL)**
 - Method for extracting rules from individual observations through an explanation.
- **Explanation**
 - Stick holds the food over the fire while keeping hands safe.
- **Generalization**
 - Any long, rigid, sharp object can be used to toast food over the fire.
 - General rule follows logically from the background knowledge of the cavemen's usual cooking process.



Chu

Natural Language Inference

T: IBM **cleared** \$18.2 billion in the first quarter.
H: IBM's **revenue** in the first quarter was \$18.2 billion.



Entailment?

YES

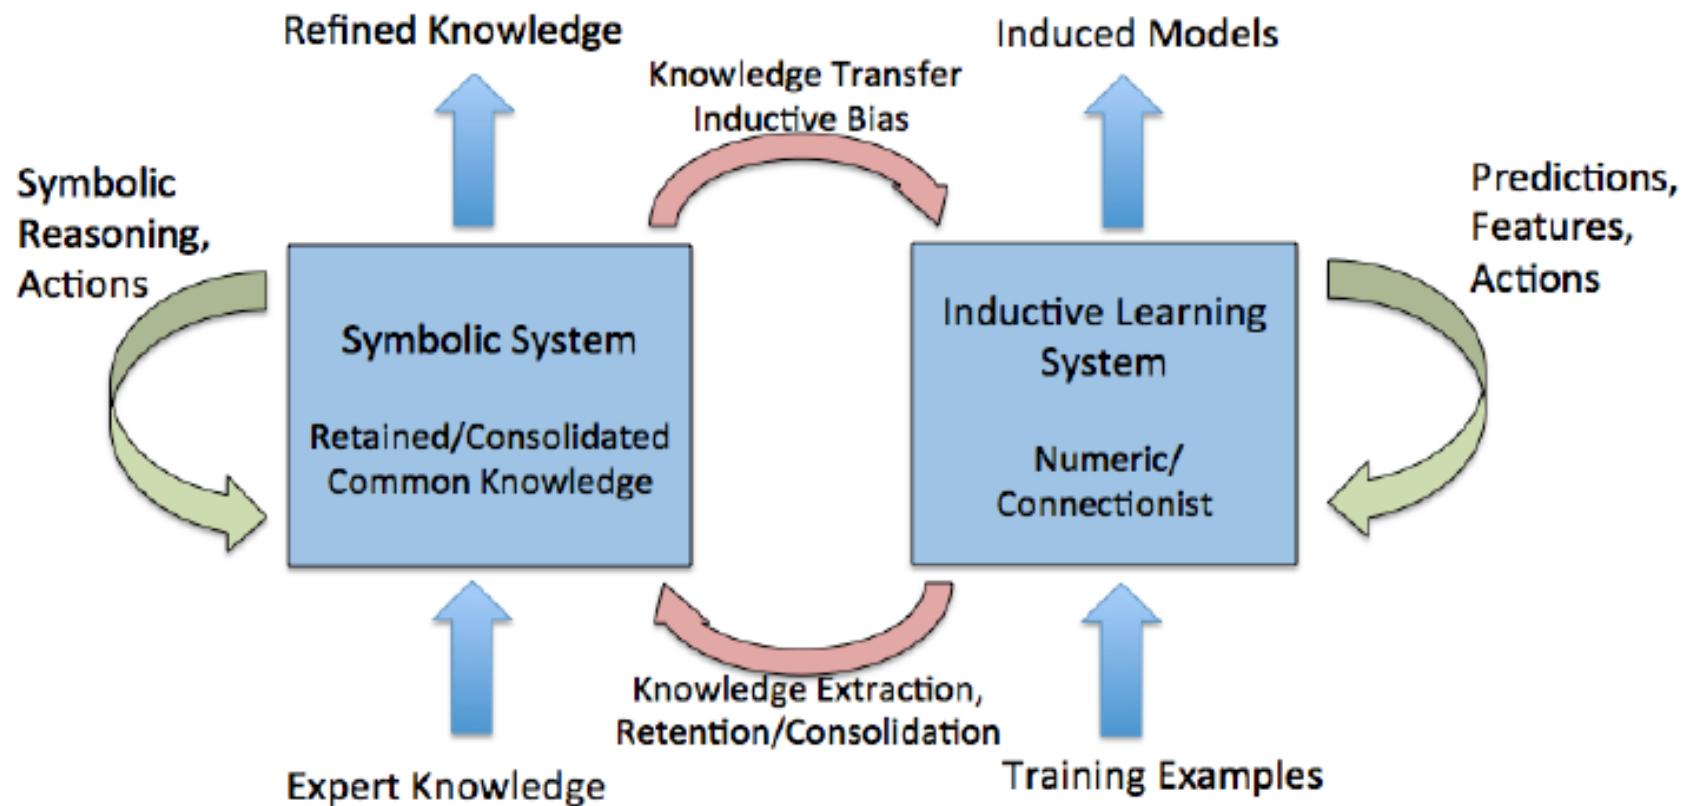
Why?

- To clear is to yield as a net profit
- A net profit is an excess of revenues over outlays in a given period of time

Natural Language Inference

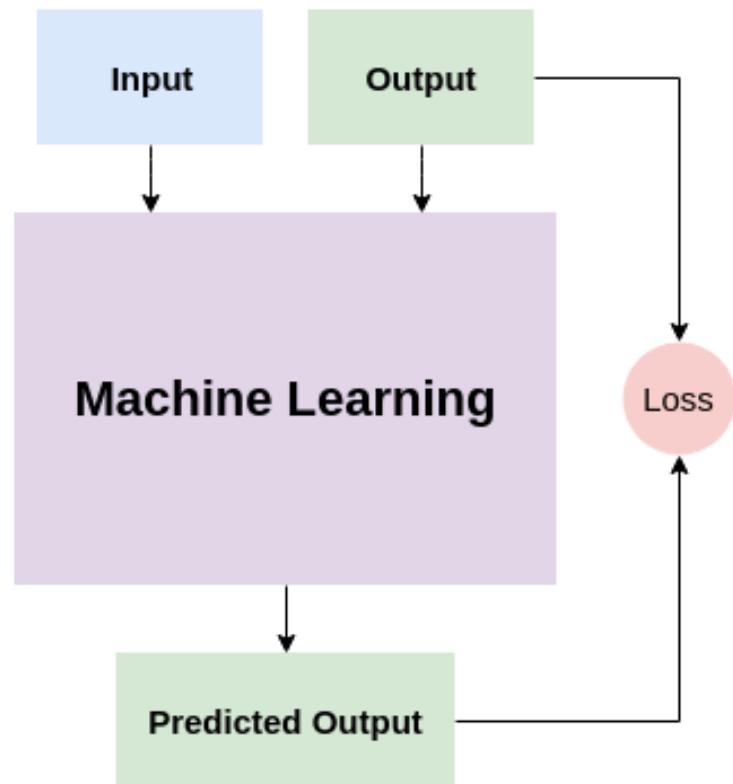
diagram	symbol	name	example
	$x \equiv y$	equivalence	<i>couch</i> \equiv <i>sofa</i>
	$x \sqsubset y$	forward entailment (strict)	<i>crow</i> \sqsubset <i>bird</i>
	$x \sqsupset y$	reverse entailment (strict)	<i>European</i> \sqsupset <i>French</i>
	$x \wedge y$	negation (exhaustive exclusion)	<i>human</i> \wedge <i>nonhuman</i>
	$x \mid y$	alternation (non-exhaustive exclusion)	<i>cat</i> \mid <i>dog</i>
	$x \square y$	cover (exhaustive non-exclusion)	<i>animal</i> \square <i>nonhuman</i>
	$x \# y$	independence	<i>hungry</i> $\#$ <i>hippo</i>

Neuro-Symbolic Models

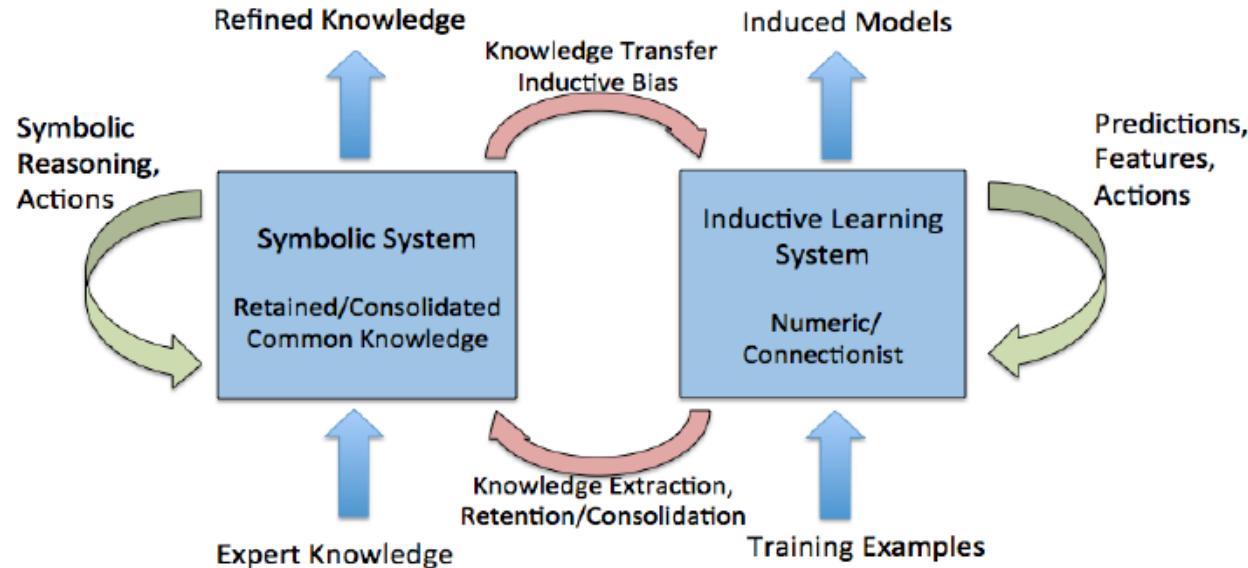


Current Limits of Deep Learning

1. Deep learning thus far is data hungry
2. Deep learning thus far is shallow and has limited capacity for transfer
3. Deep learning thus far has no natural way to deal with hierarchical structure
4. Deep learning thus far is not sufficiently transparent
5. Deep learning thus far has not been well integrated with prior knowledge
6. Deep learning thus far cannot inherently distinguish causation from correlation
7. Deep learning presumes a largely stable world, in ways that may be problematic
8. Deep learning thus far works well as an approximation, but its answers often cannot be fully trusted
9. Deep learning thus far is difficult to engineer with



Statistical vs Symbolic AI Systems



	Statistical	Symbolic
Explainability	Hard	Easy
Generalizing algebraic operations	Hard	Easy
Robustness to noise	Easy	Hard
Robustness to ambiguity	Easy	Hard
Robustness to mislabeling	Easy	Hard

∂ ILP

It is possible for systems to combine statistical perceptual with conceptual interpretable reasoning!

	Deep Learning	Symbolic Program Synthesis	∂ ILP
Robust to noise	YES	NO	YES
Can learn from non-symbolic data	YES	NO	YES
Data efficient	NO	YES	YES
Interpretable	NO	YES	YES

Evans, Richard, and Edward Grefenstette. "Learning explanatory rules from noisy data." *Journal of Artificial Intelligence Research* 61 (2018): 1-64.

Gated Graph Neural Networks

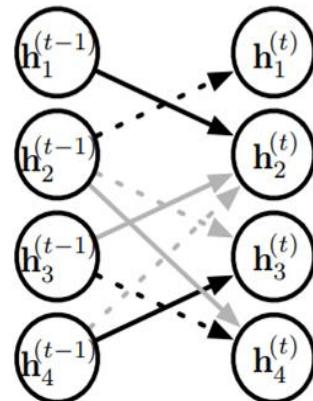
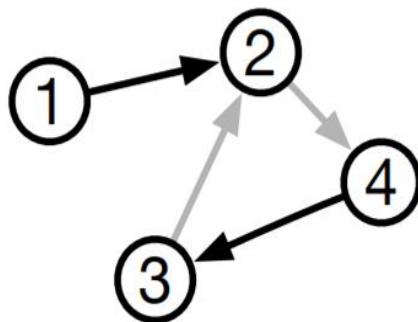
GGNNs are a neural network architecture defined according to a graph structure $G = (V, E)$

Nodes $v \in V$ take unique values from $1, \dots, |V|$, and edges are pairs $e = (v, v_0) \in V \times V$

GGNNs map graphs to outputs via two steps. First, there is a **propagation step** that computes node representations for each node; second, an **output model** $ov = g(hv, lv)$ maps from node representations and corresponding labels to an output ov for each $v \in V$

The propagation model is similar to an LSTM. Each node in the graph v has a hidden state representation $h(t)v$ that is updated at every time step t . The computation starts at $t = 0$ with initial hidden states xv that depends on the problem.

The structure of the graph, encoded in a matrix A serves to retrieve the hidden states of adjacent nodes based on the edge types between them. The hidden states are then updated by a gated update module.



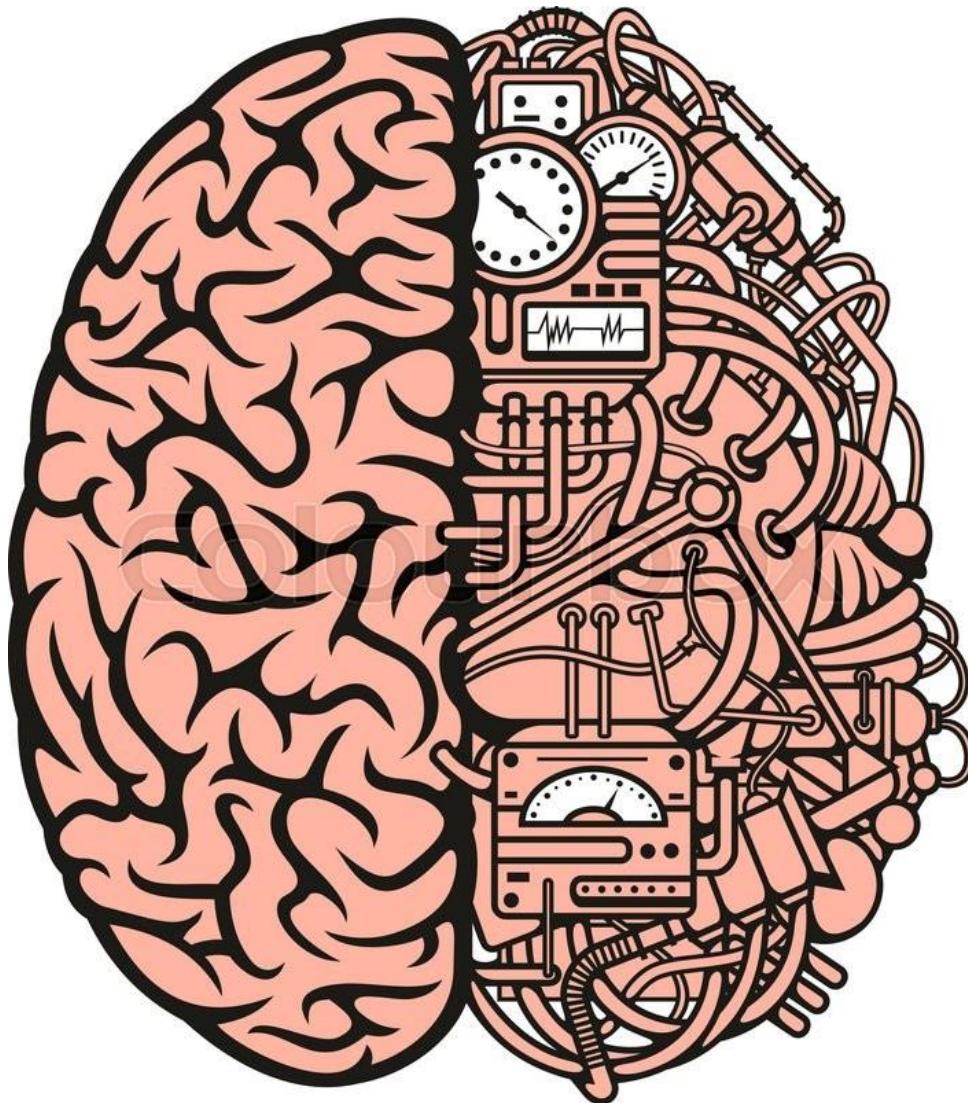
	Outgoing Edges				Incoming Edges			
	1	2	3	4	1	2	3	4
1	B							
2		C	B'				C'	
3	C							B'
4		B			C'			

(a)

(b)

(c) $\mathbf{A} = [\mathbf{A}^{(\text{out})}, \mathbf{A}^{(\text{in})}]$

Building Neuro-Symbolic Systems



Explainable QA

Q: The Pollution Prevention Act of 1990 expanded a publicly available or private database?

A: publicly available

Explanation:

Paragraph A: Pollution Prevention Act of 1990

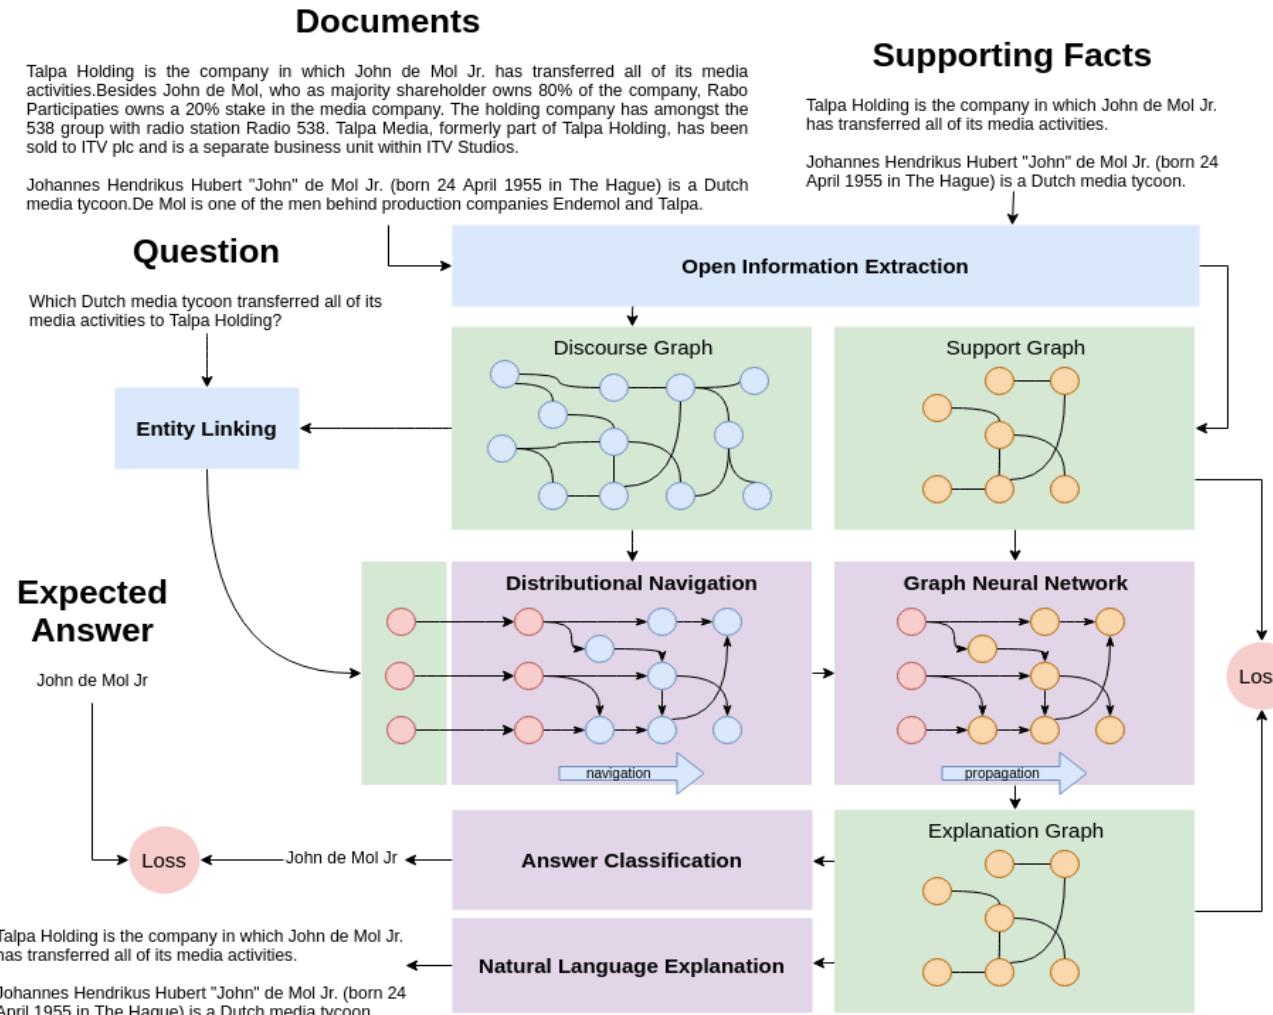
The Pollution Prevention Act of 1990 (PPA) in the United States created a national policy to have pollution prevented or reduced at the source wherever possible. It also expanded the Toxics Release Inventory.

Paragraph B: Toxics Release Inventory

The Toxics Release Inventory (TRI) is a publicly available database containing information on toxic chemical releases and other waste management activities in the United States.

Explainable QA

Neuro-Symbolic =
Knowledge Graphs +
Gated Graph Neural
Networks (GGNN)



Explainable Science QA

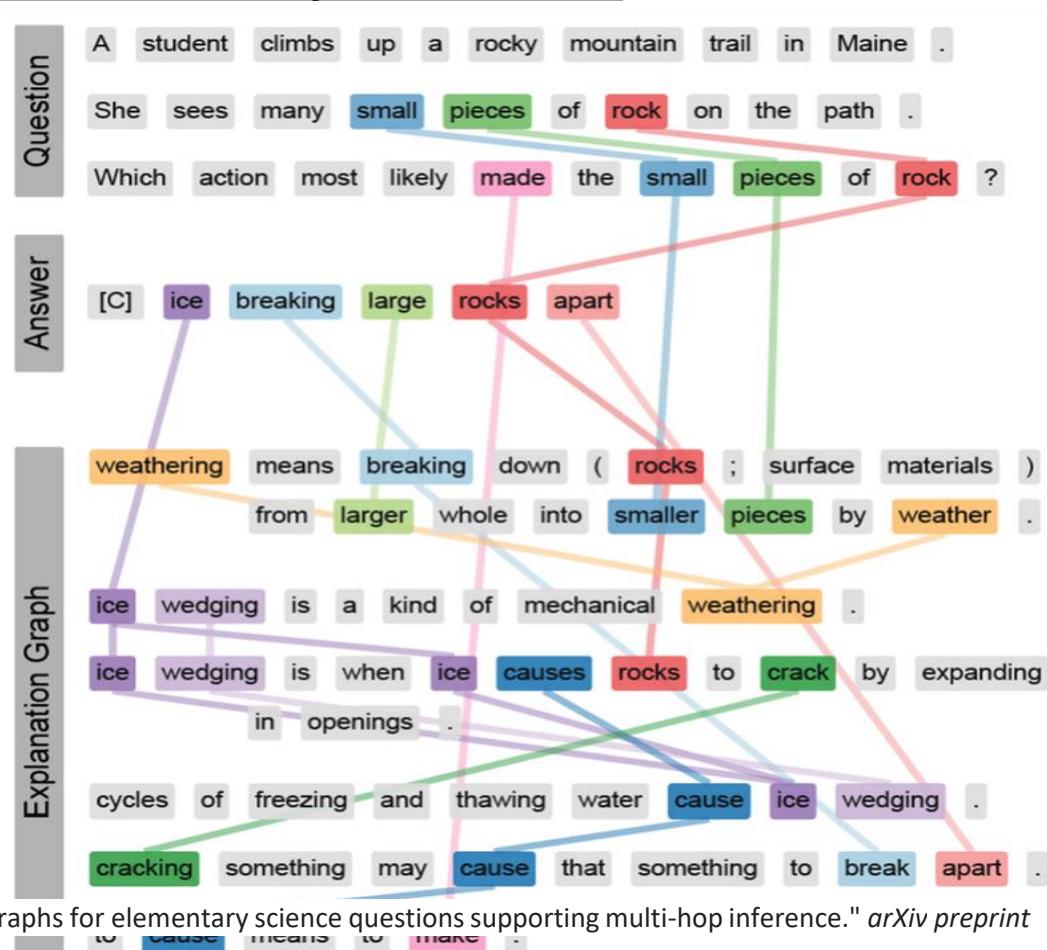
Q: A student climbs up a rocky mountain trail in Maine. She sees many small pieces of rock on the path. Which action most likely made the small pieces of rock?

- [0]: sand blowing into cracks
- [1]: leaves pressing down tightly
- [2]: ice breaking large rocks apart
- [3]: shells and bones sticking together

A: ice breaking large rocks apart

Explanation:

- weathering means breaking down (rocks) by weather
- ice wedging is a kind of mechanical weathering
- ice wedging is when ice causes rocks to crack
- cycles of freezing and thawing water cause ice wedging
- cracking something may cause that something to break
- to cause means to make





Powered by Thomson Reuters

[News Feed](#) [Connection](#) [Portfolio](#)

[Settings](#) [All](#) [Pinned](#)

3 number of pinned articles

[BRIEF-Avio receives EUR 40 mln financing from European Investment Bank <SPA2.MI>](#)

Companies: Avio SPA 80%, Avio SPA 80%, Avio SPA 80%

Topics: Business Finance, Contracts / Business Deals Events: ContactDetails

Industry: Spacecraft Manufacturing

Publication date: Oct 6, 2017 10:24:02 AM

Airbus SE

BAE Systems PLC

Boeing Co

[BRIEF-Airtelis orders three H215 Airbus Helicopters<AIR.PA><AIRG.DE>](#)

Companies: Airbus Helicopters SAS 50%, Airbus Helicopters SAS 50%, Airbus Helicopters SAS 50%

Topics: N/A Events: BusinessRelation

Industry: Aerospace & Defense - NEC, Aircraft Parts Manufacturing - NEC

Publication date: Oct 4, 2017 10:46:00 AM

Raytheon Co

BAE Systems PLC

Boeing Co

[BRIEF-British Airline Pilots' Association - pilots union calls for investigation into collapse of Monarch Airlines<MONA.UL>](#)

Companies: Monarch Airlines Ltd 80%, Monarch Airlines Ltd 80%, Monarch Airlines Ltd 80%

Topics: Other Events: N/A

Industry: Airlines - NEC

Publication date: Oct 9, 2017 1:55:13 PM

Raytheon Co

BAE Systems PLC

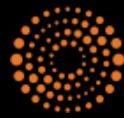
Boeing Co



THOMSON REUTERS

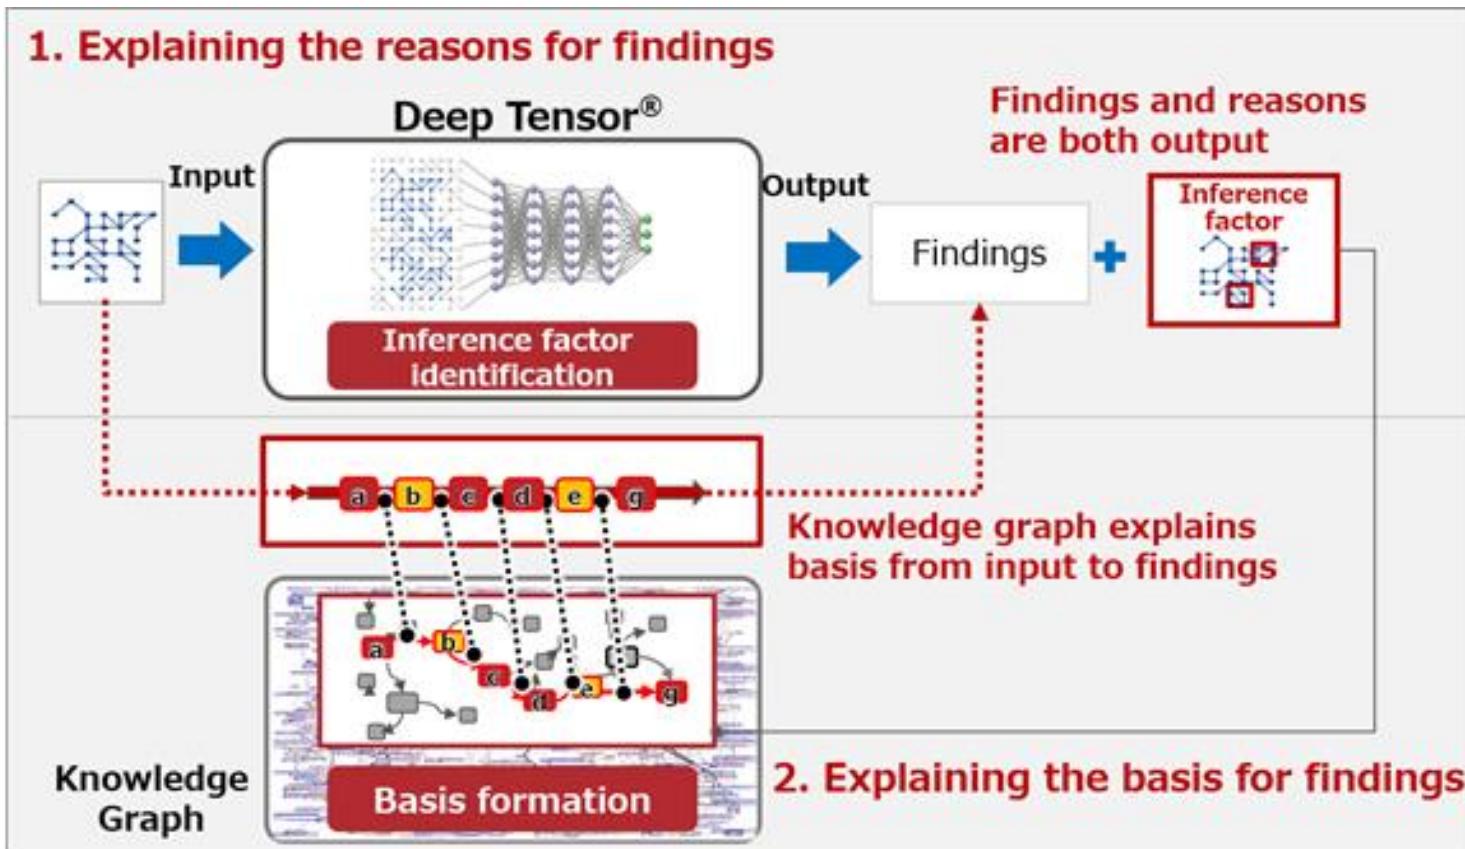
BRIEF-Avio receives EUR 40 mn financing from European Investment Bank <SPA2.MI>

Oct 6 (Reuters) - AVIO SPA <SPA2.MI> * SAYS SIGNED WITH EUROPEAN INVESTMENT BANK CONTRACT FOR EUR 40 MILLION FINANCING Source text for Eikon: [ID:nBIA5D9Ntk] Further company coverage: [SPA2.MI] (Gdynia Newsroom) (gdynia.newsroom@thomsonreuters.com; +48 58 772 0920 ;)



THOMSON REUTERS

Explainable Findings From Tensor Inferences Back to KGs



Summary of Today

- End-to-end overview of this part of the course and its underlying motivation.
- Representation, semantics, learning/inference.
- Representation and NL.
- Dialogue between Statistical and Symbolic approaches.

Next Class

- We will jump directly into Knowledge Representation (beyond FOL).
- Frames, Prototypes, Ontologies.
- Representing more complex NL discourse.

Recommended Reading

Deep Learning: A Critical Appraisal

Gary Marcus¹
New York University

<https://arxiv.org/ftp/arxiv/papers/1801/1801.00631.pdf>